

Python Team Project Report

Felipe da Paixao, Agustin Moya Rojas, Henry Pham, Aleksander Witt

LC1: Team 15

ENGR 13300

Prof. Soudebeh Taghian

Oct 23rd, 2025

1. Project Motivation

Our motivation for this project fundamentally strands from learning important skills and techniques for image processing using python, that not only teach us problem solving in a highly technical setting, but also a potentially useful skill for future applications in our work. This is because several engineering disciplines use image processing for crucial tasks. In mechanical engineering [Henry], very detailed image processing is used for automated quality control systems and robotics in the automotive industry (*Gardel, Science Direct*), in the electrical engineering [Aleks] field it is increasingly being used, paired with trained AI models, to visually inspect client's PCBs and give feedback on how to optimize them (*Averroes*), and in IBE [Agustin & Felipe] as it has varying relevance depending on the specific engineering discipline we end up choosing, a direct relation to business, among many others, could be the growing applications of "shelf intelligence" to optimize warehouse and inventory systems in real world retail businesses (*Trax*). It is evident that all these are direct applications of our current project for future applications we may encounter in our daily work or projects, and even if the scope varies, the experience of problem solving by trial and error and growing our teamwork skills, is of great value to us.

2. Project Overview and Methods

In this project we were tasked to develop a program for image processing and created algorithms for each step of the process including preprocessing and image extraction and finally using this extracted data along with machine learning algorithms to predict accurate and reliable new data as a result. In checkpoint 1, it was crucial to first correctly prepare the images, before extracting its features. To do this, we make sure that the image properties are consistent and correctly validated and convert them to the format of our choice (HSV, RGB, Grayscale), resize them, and finally have methods in place to handle errors.


Next, for checkpoint 2 we moved on to the more practical and applicable steps, to actually extract the image's shape-based features, for example; edge detection and line detection. We used filtering techniques like a gaussian filter to make images less complex and easier to extract features from and then a Sobel filter to determine where edges were in the image based on the brightness of pixels. From there, we could then convert the images to HSV and calculate statistical values like the mean and standard deviation for the image's hue, saturation, and value. We also detected shapes by counting lines and detecting circles to give us data that could be used to classify the images.

Finally, in checkpoint 3, we used all the features extracted from previous tasks to train machine learning algorithms to classify these images. Our KNN model compared each image to others in the dataset to classify it based on the images that are most similar. The Decision Tree model splits the data into left and right branches based on feature values and whether the data met the thresholds of the feature values to predict the image classification. The Linear Regression model calculated a best fit line or curve that separated classes to determine the probability of an image fitting in that class. Our data was split into training, validation, and test sets to evaluate the model's performance.

3. Discussion of Algorithm Design


Checkpoint 1 Flowcharts:

[tp1_team_1_15.pdf](#)




For this task, we designed a function to load images. To do so, we made a user defined function utilizing the Pillow library to load an image, convert it to RGB, then convert it to an array to work with. From there, we normalized the image and optionally converted it to greyscale. Doing so provided us with a modular and easy to use interface for loading images and converting them to arrays, which we heavily relied on later.

tp1_team_2_15.pdf



For this task, we designed a function that takes an image and then cleans and pads it. Doing so standardizes the arrays, preventing unwanted noise, outliers, or size differences between images. Doing this allows our later analysis to be more accurate, as all imported images share a size and are clean. Furthermore, defining this in a user defined function makes it easy to use later on.

tp1_team_3_15.pdf




In this assignment, we created two functions in order to convert a whole image into HSV from RGB.

First, we designed a function that iterates over a single pixel, determines the RGB values, then converts to HSV for that single pixel. From there, we created another function that iterates over every pixel in a cleaned and padded image, using the previous function to convert the entire image to HSV. By doing so, we make color detection much easier, as HSV images are far easier to detect changes in color than RGB ones.


Checkpoint 2 Flowcharts:

tp2_team_1_15.pdf



In this task, we designed a function to apply a gaussian filter to a greyscale image. To do so, we used the functions from previous tasks to load an image, pad it, and convert it to grayscale. From there, we defined a Gaussian filter kernel which we then applied to every pixel in the image, returning an array of a filtered image. Doing this blurs the image, making later detection of edges far easier and reducing noise. Therefore, the Gaussian filter is a vital part of many image processing applications.


tp2_team_2_15.pdf



For this task, we designed a function that applies a Sobel filter to the image processed with a Gaussian process prior. To do so, we initialized a kernel for both the x and y gradients in the image, then used those values to find the Sobel filtered array for every pixel in a 3x3 area of each iterated pixel, eventually producing a new image with only the edges visible. Doing so allows us to easily analyze images based on the edges, as the function clearly defines them without any noise. Because of this fact, Sobel filters are necessary when trying to analyze the edges of multiple images.


tp2_team_3_15.pdf

Unfortunately, our team was unable to complete task three. Because of this, we later used the provided CSV for our data. The following flow chart explains how the extract features function, which we got working works.




Checkpoint 3 Flowcharts:

tp3_team_1_15.pdf



For this assignment, we built a KNN (K-Nearest Classifier) in order to use our model and detect stop signs. To do so, we made functions that loaded a dataset, trained the dataset, and predicted the results in order to analyze our data. In doing so, we received promising results. We received a training accuracy of 91 percent, which meant that our model was successful. However, this method of classification was not as accurate as the logistic regression seen below.


tp3_team_2_15.pdf



For this assignment, we used a decision tree classifier to analyze our data. Simply put, we defined functions to calculate the entropy of our data, the information gain from each split, and building the decision tree itself. The specifics of our design can be seen in the flowchart above. Ultimately, this

method gave us similar results to the KNN, just with a different method. This exposed us to the fact that there are many ways to solve the same problem within image processing.

tp3_team_3_15.pdf



Finally, we created a logistic regression program in order to analyze our data in a third way. This time, we calculated the probability of each value occurring, predicted the labels, calculated our loss and gradients, then used this data to train a model and compare results. Once again, this model gave us similar results to the other two. While similar, receiving similar results with different methods of classification revealed that there are many approaches to solving the same problem, simply using different methods for different use cases.

4. References

- Gardel, A., Molina, J., Garcia, I., and Bravo, J.L. "Applications of Image Processing in Robotics and Instrumentation." *Mechanical Systems and Signal Processing*, vol. 124, 2019, pp. 142-169. Elsevier/ScienceDirect. <https://doi.org/10.1016/j.ymssp.2018.05.048>
- Averroes AI. "AI Visual Inspection for PCB & Electronics Manufacturing." *Averroes Technologies*, 2024. <https://averroes.ai/industry/ai-electronics>. Accessed 23 Oct. 2025.
- Trax Retail. "Revolutionizing Retail Through AI-Driven Data and Consumer Engagement." *Trax Retail*, 2025. <https://traxretail.com/>. Accessed 23 Oct. 2025.

5. Appendix

1. User manual

Checkpoint 1:

- ***Task 1-input an image to load and confirm if you want it to be converted to grayscale***



-

```
PS C:\Users\henry\OneDrive\ENGR_133\ENGR 133\TP_1> & C:/Users/henry/AppData/Local/Microsoft/WindowsApps/python3.11.exe
  "c:/Users/henry/OneDrive/ENGR_133/ENGR 133/TP_1/tp1_team_1_15.py"
Enter the path of the image you want to load: ref_col.png
Would you like to convert to grayscale?
yes
PS C:\Users\henry\OneDrive\ENGR_133\ENGR 133\TP_1>
```

- ***Task 2 – input an image to clean and resize***



```
Enter the path of the image you want to clean: ref_col_raw0.png
Image shape before cleaning: (112, 121, 3)
Resized image to: (92, 100)
• Image shape after cleaning: (100, 100, 3) □
```

- ***Task 3 – input an RGB image to be converted to HSV and input XY coordinates of a pixel to be inspected***



```
Enter the path of the RGB image you want to convert to hsv: ref_col_raw4.png
Resized image to: (92, 100)
Enter the x and y coordinates of the pixel you want inspect: 50, 50
RGB values of the (50, 50) pixel: R=255, G=255, B=255
Converting ref_col_raw4.png to HSV...
• HSV values of the (50, 50) pixel: H=0, S=0, V=255
```


Checkpoint 2:

- *Task 1 – input an image to be blurred using a gaussian filter*




- Enter the path to the image file: `ref_col.png`

- *Task 2 – input an image to define its edges using a sobel filter*



- Enter the path to the image file: `ref_col.png`

- *Task 3 – input a dataset folder, metadata file, csv feature file, and x and y axis features*



```

Enter the name of the dataset folder:
ML_Images
Enter the name of the metadata file:
img_metadata.csv
Enter the name of output csv features file:
img_features.csv

Training Metadata DataFrame:
   Width  Height ClassId      Path
0     26      29      1 1_00015_00000.png
1     27      30      1 1_00015_00001.png
2     27      30      1 1_00015_00002.png
3     28      31      1 1_00015_00003.png
4     28      30      1 1_00015_00004.png

Features have been extracted and saved to img_features.csv

Feature Dataset Shape:
(1469, 10)

Feature Dataset Head:
   hue_mean  hue_std  saturation_mean  saturation_std  value_mean  value_std  num_lines  has_circle      Path  ClassId
0    77.5901  74.895838        87.5351    76.634594    16.0337  17.254430      0.0       0.0 1_00015_00000.png      1
1    71.8692  73.134396        91.1982    74.631321    16.4933  17.784649      0.0       0.0 1_00015_00001.png      1
2    78.3621  73.230186        91.6578    76.665620    16.7615  17.559311      0.0       0.0 1_00015_00002.png      1
3    77.3715  71.860282        92.6498    75.498783    16.5275  17.430635      0.0       0.0 1_00015_00003.png      1
4    81.7215  70.873320       101.1194    77.180979    17.8706  18.428946      0.0       0.0 1_00015_00004.png      1

Enter the column to use as the x-axis feature:
hue_mean


Enter the column to use as the y-axis feature:
saturation_mean

● Scatter plot saved to KNN_feature_space.png

```

Checkpoint 3:

- **Task 1 – input csv file, whether to shuffle, and seed number**



```
Enter the path to the feature dataset: img_features.csv
Shuffle the dataset? (yes/no): no
Enter a seed for loading the dataset: 70
```

```
Data loaded and split into
```

```
Training set:
size: 1175
Validation set:
size: 147
Test set:
size: 147
```

```
Based on the plots, the best k appears to be: 200
```

```
Evaluating final model on test set with k = 200...
```


```
--- Final Model Performance ---
```

```
Test Set Accuracy: 0.6667
```

```
Test Set Error Rate: 0.3333
```

- *Task 2 – input feature dataset, optimal max depth, and optimal min samples*

ENGR 133: Transforming Ideas to Innovation



```
Enter the path to the feature dataset: img_features.csv

Data loaded and split:
Training set: size: 1175
Validation set: size: 147
Test set: size: 147

Finding optimal tree max depth and min samples split...


Enter the optimal max depth according to your interpretation of the plot:
3

Enter the optimal min samples split according to your interpretation of the plot:
30
Evaluating final model on test set with the chosen hyperparameters

--- Final Model Performance ---
Test Set Accuracy: 0.8912
Validation Set Accuracy: 0.8912
Training Set Accuracy: 0.9021

--- Final Decision Tree Rules ---
If hue_mean < 61.62:
    If hue_mean < 58.24:
        Predict: Stop Sign
    Else (hue_mean > 58.24)
        If value_mean < 43.09:
            Predict: Stop Sign
        Else (value_mean > 43.09)
            Predict: Stop Sign
    Else (hue_mean > 61.62)
        If hue_std < 82.89:
            If hue_mean < 70.80:
                Predict: Right Turn Sign
            Else (hue_mean > 70.80)
                Predict: Right Turn Sign
        Else (hue_std > 82.89)
            If num_lines < 15.50:
                Predict: Right Turn Sign
            Else (num_lines > 15.50)
                Predict: Stop Sign
•
```

- **Task 3 – outputs logistic regression model**



```
Iteration 0, Loss: 0.6916
Iteration 100, Loss: 0.3735
Iteration 200, Loss: 0.3314
Iteration 300, Loss: 0.3129
Iteration 400, Loss: 0.3030
Iteration 500, Loss: 0.2973
Iteration 600, Loss: 0.2936
Iteration 700, Loss: 0.2912
Iteration 800, Loss: 0.2895
Iteration 900, Loss: 0.2883
Iteration 1000, Loss: 0.2875
Iteration 1100, Loss: 0.2869
Iteration 1200, Loss: 0.2864
Iteration 1300, Loss: 0.2860
Iteration 1400, Loss: 0.2858
Iteration 1500, Loss: 0.2856
Iteration 1600, Loss: 0.2854
Iteration 1700, Loss: 0.2853
Iteration 1800, Loss: 0.2852
Iteration 1900, Loss: 0.2851
Validation accuracy: 0.9048
Test accuracy 0.9184
Validation error 0.0952
• Test error 0.0816
```

2. Project management plan (Team members' contribution)

Our methods for collaboration during this team project included organizing the tasks we needed to get done by the weekly deadlines and continuously communicating in our group chat to keep us on track, and check if anyone needed support. By working on our tasks together, we ultimately supported each other in every step of the way and learned from each other's skills and insights as we worked. As we worked on tasks together, this involved delegating certain small components of the task between members and continuously checking on their progress and understanding of it, to then bring all components together.

The general contribution of each team member was the following; Henry focused on debugging and communicating task progression with the team. He led the coding on the

decision tree model and the gaussian filter implementation. One improvement for Henry would be to add more thorough commenting on his code to help others understand its functions. Agustin led the coding for Checkpoint 1 Task 1 along with Checkpoint 3 Task 1 and worked closely to answer project report questions. One improvement for Agustin would be to be more detailed when creating flowcharts to convey the right processes to users. Aleks led the coding for Checkpoint 1 Task 3, Checkpoint 2 Task 2, and Checkpoint 3 Task 3. Additionally, Aleks significantly contributed to checkpoint 3 task 1 after errors with our code. One improvement for Aleks would be to write cleaner and more standard code, allowing team members to read it easier. Felipe led the coding for Checkpoint 2 Task 3, helped with Checkpoint 3 Task 1, and worked closely with team members on other smaller pieces of code in other tasks. One improvement for Felipe would be to work farther away from the deadline to have enough time to debug and get help from teammates if needed.

3. Discussion of design process (*Approach to the design process*)

Our design process fundamentally involves working together to assess the general project, idea, and specific tasks assigned, and then following a design-oriented process to complete this project. This process first included talking about our approach and how we would tackle each task, and how do the tasks complement each other, as in both checkpoint 1 and 2, many of the latter tasks like #2, #3, included using functions completed in task #1, and continued to improve on them, or modify their usage. This is why our approach had to be sequential or progressive, rather than recur to delegation and separation of tasks. After this, it was important to create flowcharts for each task, to have a conceptual understanding of what we had to do, and the functions we had to use, and after deliberation on our approach, we got to work. This part was the bulk of the project, coding our specific

functions and main functions, and after we had a rough working draft, we also stepped back and thought if this was the intended way to do it, or if there was an easier, more efficient way too.

4. **Code**

tp1_team_1_15.py

```
import numpy as np

from PIL import Image

import matplotlib.pyplot as plt


# load image

def load_img(path):

    try: #Try function to check for errors; looked up method to check for errors.

        img = Image.open(path)

    except FileNotFoundError: #Says that if a file not found error is detected to print an error.

        print(f"error {path} not found.")

    return None


# remove 4th channel

if img.mode == 'RGBA':


    img=img.convert('RGB')
```

```
#convert to array

pixel_array = np.array(img)

# normalize

normalized_array = pixel_array/255.0

# Linearizing the Image

img_arr = np.where(normalized_array <= 0.04045, normalized_array /
12.92, ((normalized_array + 0.055) / 1.055) ** 2.4)

#unnormalizes and makes sure it's 8bit data

img_array=(img_arr*255).astype(np.uint8)

return img_array

# grayscale conversion

def rgb_to_grayscale(img_array):

    # check for color image

    if img_array.ndim != 3:

        print("Error")

    return img_array

#gets color channels

red= img_array[:, :, 0]
```

```
green= img_array[:, :, 1]

blue = img_array[:, :, 2]

#appends array of gray values

gray_arr = (0.2126 * red) + (0.7152 * green) + (0.0722 * blue)

#makes sure it's 8bit

gray_array=(gray_arr).astype(np.uint8)

return gray_array

# main

def main():

    path = input("Enter the path of the image you want to load: ")

    loaded_img = load_img(path)

    final_img = loaded_img

    # if color, convert?

    if loaded_img.ndim == 3: #so a color one will have 3 channels

        choice = input("Would you like to convert to grayscale?\n")

        if choice == 'yes':

            final_img = rgb_to_grayscale(loaded_img)
```

```
#cmap for gray

if final_img.ndim == 2: # so gray

    plt.imshow(final_img, cmap='gray')

else:

    plt.imshow(final_img)

    plt.axis('off')

plt.show()

if __name__ == "__main__":
    main()
```

tp1_team_2_15.py

```
from PIL import Image, ImageOps

import numpy as np

from tp1_team_1_15 import load_img

import matplotlib.pyplot as plt

def clean_image(arr):

    img = Image.fromarray(arr) #convert array back to image
```

```
width, height = img.size

aspect_ratio = width/height

#resize image based on aspect ratio

if aspect_ratio > 1: #wider than tall

    new_w = 100

    new_h = int(100/aspect_ratio)

elif aspect_ratio < 1: #taller than wide

    new_w = int(100*aspect_ratio)

    new_h = 100

else: #square

    new_w = 100

    new_h = 100

img = img.resize((new_w, new_h), Image.BILINEAR) #resizes image

print(f'Resized image to: {(img.size[1], img.size[0])}') #lists height then width to match samples

if img.mode == "RGB": #creates fill color for color images or grayscale images

    fill_color = (0, 0, 0)

else:
```

```
fill_color = 0

img_padded = ImageOps.pad(img, (100, 100), method = Image.BILINEAR, color = fill_color)
#ImageOps.pad pads the image to the given dimensions with given color

new_img = np.array(img_padded) #convert back to array

return new_img #returns cleaned array

def main():

    path = input("Enter the path of the image you want to clean: ") #takes input

    img_array = load_img(path)

    print(f"Image shape before cleaning: {img_array.shape}") #dimensions of array

    cleaned_img = clean_image(img_array) #cleans image

    print(f"Image shape after cleaning: {cleaned_img.shape}")

    if cleaned_img.ndim == 3: #display image for color or grayscale

        plt.imshow(cleaned_img)

    else:

        plt.imshow(cleaned_img, cmap = "gray")

    plt.axis('off')

    plt.show()
```

```
if __name__ == "__main__":
```

```
    main()
```

tp1_team_3_15.py

```
from pathlib import Path
```

```
from PIL import Image # Importing libraries
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from tp1_team_1_15 import load_img
```

```
from tp1_team_2_15 import clean_image
```

```
def rgb_to_hsv(r,g,b):
```

```
r2=r/255.0 #Added .0 to every number to try to fix floating point differences.
```

```
b2=b/255.0 #Normalizing values
```

```
g2=g/255.0
```

```
cmax=max(r2,g2,b2) #Finding the max and minimum RGB value and their difference
```

```
cmin=min(r2,g2,b2)
```

```
difference=cmax-cmin
```

```
if difference == 0.0:
```

```
H2=0.0

if cmax==r2: #Defining the Hue based on the most prominent RGB value

    H2=(60.0*((g2-b2)/difference))%360.0 #Added 360 to account for negatives.

elif cmax==g2:

    H2=((60.0*((b2-r2)/difference))+120.0)%360.0

elif cmax==b2:

    H2=((60.0*((r2-g2)/difference))+240.0)%360.0

if cmax == 0:

    S2=0.0 #Defining Saturation (0 if black)

else:

    S2 = (difference/cmax)

V2=cmax #The value is just the max rgb value

H=(H2/360.0)*255.0 #Not entirely sure why H2 must be divided by 360 but the assignment says so.

S=S2*255.0 #Scaling the values back up.

V=V2*255.0

return H,S,V

def convert_to_hsv(rgb_image):
```

```
    hsv_im=rgb_image.astype(float).copy() #Copying the rgb array to a new array to work on;  
    converting to float from unassigned integer.  
  
    shape=hsv_im.shape #Finding the shape  
  
    h=shape[0]  
  
    w=shape[1]  
  
    for i in range(h): #For loop to change every pixel (done in previous assignments)  
  
        for k in range(w):  
  
            r=rgb_image[i,k,0]  
  
            g=rgb_image[i,k,1]  
  
            b=rgb_image[i,k,2]  
  
            H,S,V=rgb_to_hsv(r,g,b) #Actually converting  
  
            hsv_im[i,k,0]=H  
  
            hsv_im[i,k,1]=S #Actually converting each one to HSV  
  
            hsv_im[i,k,2]=V  
  
    hsv_image=(hsv_im).astype(np.uint8) #Converting back to uint8  
  
    return hsv_image
```

```
def main():

    path = input("Enter the path of the RGB image you want to convert to hsv: ") #Defining file

    arr=load_img(path)

    clean_img=clean_image(arr) #Running the udfs

    inp=(input("Enter the x and y coordinates of the pixel you want inspect: "))

    x,y=inp.split(',') #Splitting the input to seperate x and y. Did this before in the assignment where
    we had to split up the arrays. Pretask something i'm pretty sure.

    R,G,B=clean_img[int(x),int(y)].astype(int)

    print(f'RGB values of the ({x},{y}) pixel: R={R}, G={G}, B={B}')

    print(f'Converting {path} to HSV...') #All self explanatory.

    H,S,V=rgb_to_hsv(R,G,B)

    print(f'HSV values of the ({x},{y}) pixel: H={int(H)}, S={int(S)}, V={int(V)}) #Used int instead of
    .0f to get autograder to work.

    hsv_img=convert_to_hsv(clean_img) #Converting to HSV

    plt.imshow(hsv_img)

    plt.axis('off')

    plt.show()
```

```
if __name__ == "__main__":
    main()

tp2_team_1_15.py
import numpy as np

from PIL import Image

import math

from tp1_team_1_15 import load_img

from tp1_team_1_15 import rgb_to_grayscale

import matplotlib.pyplot as plt

def gaussian_filter(gry_arr, std_dev):
    gry_arr = gry_arr.astype(np.float64)

    kernel_size = int(2 * math.ceil(3 * std_dev) + 1) #calculate kernel size

    pad_size = int((kernel_size - 1)/2) #calculate size to pad image by

    x, y = np.meshgrid(np.arange(-(kernel_size//2), kernel_size//2 + 1), np.arange(-(kernel_size//2), kernel_size//2 + 1)) #creates grid for x and y values

    kernel = (np.exp(-(x**2+y**2)/(2 * std_dev**2))) #uses gaussian function that gives weights for each pixel in kernel
```

```
kernel = kernel/np.sum(kernel, dtype = np.float64) #normalize values to be within 0, 1

padded_arr = np.pad(gry_arr, pad_size, mode = "constant")#creates padded array with zeros

height, width = gry_arr.shape

output = np.zeros((height, width), dtype = float) #creates blank output array

#iterates through padded array

for r in range(height):

    for c in range(width):

        region = padded_arr[r:r + kernel_size, c:c + kernel_size] #gets 7x7 region of pixels

        output[r, c] = np.sum(region * kernel) #applies weighted pixel to output array

return (np.clip(output, 0, 255)).astype(np.uint8) #makes sure it's 8bit

def main():

    path = input("Enter the path to the image file: ")

    img = load_img(path)

    gry_arr = rgb_to_grayscale(img)
```

```
std_dev = 1

filtered_arr = gaussian_filter(gry_arr, std_dev) #blurs image

plt.imshow(filtered_arr, cmap = "gray") #show image

plt.axis('off')

plt.show()

if __name__ == "__main__":
    main()
```

tp2_team_2_15.py

```
from pathlib import Path

from PIL import Image # Importing libraries

import numpy as np

from tp1_team_1_15 import load_img

from tp1_team_1_15 import rgb_to_grayscale #Importing all the files this thing needs

import matplotlib.pyplot as plt
```

```
from tp2_team_1_15 import gaussian_filter

def sobel_filter(gaussian_filter):

    blur_pad=np.pad(gaussian_filter, (1,1)) #Using the pad function to pad the filtered image.

    kernelx=np.array([[-1,0,1],[-2,0,2],[-1,0,1]]) #Defining the two kernels

    kernaly=np.array([[-1,-2,-1],[0,0,0],[1,2,1]])

    shape=gaussian_filter.shape

    h=shape[0] #Getting the shape of the array and using that to define the height and width

    w=shape[1]

    i=0 #Setting the indices to 0

    k=0

    gx=np.zeros((h,w)) #Forming two zero arrays so I don't have to overwrite existing arrays

    gy=np.zeros((h,w))

    for i in range(h):

        for k in range(w): #For loops that go over every pixel in the image

            sum1=0 #Reseting sums for each pixel

            sum2=0

            for m in range(3):
```

```
for l in range(3): #Setting another loop that goes over kernels

    sum1+=blur_pad[i+m,k+l]*kernalx[m,l] #Multiplying each pixel in a 3x3 area of the
target pixel by the kernel

    sum2+=blur_pad[i+m,k+l]*kernaly[m,l] #Same here

    gx[i,k]=sum1 #Summing up all the values to find the gradient

    gy[i,k]=sum2

    gmag=np.sqrt(gx**2+gy**2) #Self explanatory - gives the magnitude of the gradient at each
pixel.

threshold=50 #Manual threshold setting

for i in range(h):

    for k in range(w): #Loop that checks every pixel

        if gmag[i,k]>=threshold:

            gmag[i,k]=255 #Loop that says that if gmag is greater than the threshold, the pixel has
max brightness. Otherwise it is black.

        else:

            gmag[i,k]=0

return gmag.astype(np.uint8) #Returning the sobel filtered image.
```

```
def main():

    path = input("Enter the path to the image file: ") #Self explanatory.

    img = load_img(path)

    gry_arr = rgb_to_grayscale(img)

    std_dev = 1

    filtered_arr = gaussian_filter(gry_arr, std_dev) #All of this has been done in other parts.

    sobel=sobel_filter(filtered_arr)

    plt.imshow(sobel, cmap="gray")

    plt.axis('off')

    plt.show()

if __name__ == "__main__":
    main()
```

tp2_team_3_15.py

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt
```

```
import cv2

from tp1_team_1_15 import load_img

from tp1_team_1_15 import rgb_to_grayscale

from tp1_team_2_15 import clean_image

from tp1_team_3_15 import rgb_to_hsv

from tp1_team_3_15 import convert_to_hsv

from tp2_team_1_15 import gaussian_filter

from tp2_team_2_15 import sobel_filter

# extract features function (HSV mean and sd)

def extract_features(new_img):

    # converts the input RGB image to the HSV color space

    new_hsv = convert_to_hsv(new_img)

    # calculates the mean and standard deviation for each of the three channels of the HSV image

    hue_mean = np.mean(new_hsv[:, :, 0])

    hue_std = np.std(new_hsv[:, :, 0])

    saturation_mean = np.mean(new_hsv[:, :, 1])
```

```
saturation_std = np.std(new_hsv[:, :, 1])  
  
value_mean = np.mean(new_hsv[:, :, 2])  
  
value_std = np.std(new_hsv[:, :, 2])  
  
# convert the rgb image to grayscale  
  
gray_img = rgb_to_grayscale(new_img)  
  
# Detect circles  
# Detects if a large circle is present. Returns 1 if found, 0 otherwise.  
def detect_circles(gray_img):  
  
    """Detects if a large circle is present. Returns 1 if found, 0 otherwise."""  
  
    # Hough Circles works best on a grayscale, slightly blurred image  
  
    # Apply a Gaussian blur  
  
    blurred_img = np.array(gaussian_filter(gray_img, sigma=1.5))  
  
    # Detect circles  
    circles = cv2.HoughCircles(  
        blurred_img,  
        cv2.HOUGH_GRADIENT,  
        dp=1.2, # Inverse ratio of accumulator resolution
```

```
minDist=100, # Minimum distance between centers of detected circles

param1=150, # Upper threshold for the internal Canny edge detector

param2=50, # Threshold for center detection

minRadius=20, # Minimum circle radius to detect

maxRadius=50 # Maximum circle radius to detect

)

return 1 if circles is not None else 0

def count_lines(gray_img):

    """Counts the number of lines in an image using Hough Line Transform."""

# Edge detection via sobel filtering is a prerequisite for Hough Lines

edges = sobel_filter(gray_img)

# Detect lines using the edge map

lines = cv2.HoughLinesP(edges, 1, np.pi / 180, threshold=50, minLineLength=30,
maxLineGap=10)

return len(lines) if lines is not None else 0
```

```
# runs UDFs to get the number of lines and detect the shape of the grayscale image

number_lines = count_lines(gray_img)

has_circles = detect_circles(gray_img)

return hue_mean, hue_std, saturation_mean, saturation_std, value_mean, value_std,
number_lines, has_circles

def main():

    # get inputs from the user

    folder = input("Enter the name of the dataset folder: ")

    name = input("Enter the name of the metadata file: ")

    output = input("Enter the name of output csv features file: ")

    # load the metadata csv -- used Claude AI

    metadata_path = f"{folder}/{name}"

try:

    df = pd.read_csv(metadata_path)
```



```
# extracts the 8 features from the cleaned image

hue_mean, hue_std, sat_mean, sat_std, val_mean, val_std, num_lines, has_circles =
extract_features(cleaned_img)

# returns the 10 values (2 from the original row and 8 extracted features)

return pd.Series({


    'Path': row['Path'],


    'ClassId': row['ClassId'],


    'hue_mean': hue_mean,


    'hue_std': hue_std,


    'saturation_mean': sat_mean,


    'saturation_std': sat_std,


    'value_mean': val_mean,


    'value_std': val_std,


    'number_lines': num_lines,


    'has_circles': has_circles


})
```

```
# applies the process row function to each row

features_df = df.apply(process_row, axis=1)

# saves to csv

features_df.to_csv(output, index=False)

print(f"Features have been extracted and saved to {output}")

# prints the shape and head of final dataset

print("\nFinal Dataset Shape:", features_df.shape)

print("\nFinal Dataset Head:")

print(features_df.head())

print()

# initialize variables for scatter plot from user input

#x = input("Enter the column to use as the x-axis feature: ")
```

```
#y = input("Enter the column to use as the y-axis feature: ")

# creates the scatter plot

#plt.scatter(x,y, color='blue', marker='o', s=100, alpha=0.7, edgecolors='black')

#plt.scatter(x,y, color='blue', marker='o', s=100, alpha=0.7, edgecolors='black')

# adds title and axis labels

#plt.title(f"2D Feature Space: {x} vs. {y}")

#plt.xlabel(f"{x}")

#plt.ylabel(f"{y}")

#plt.legend()

#plt.grid(True)

#plt.show

if __name__ == "__main__":
    main()
```

tp3_team_1_15.py

```
from pathlib import Path

from PIL import Image # Importing libraries

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

def load_dataset(file_path, feature_cols, label_col, shuffle, seed=70):

    df = pd.read_csv(file_path)

    X = df[feature_cols].to_numpy()

    y = df[label_col].to_numpy()

    if shuffle:

        np.random.seed(seed)

        shuffled_indices = np.random.permutation(len(y)) #This function is taken directly from the examples

        X = X[shuffled_indices]

        y = y[shuffled_indices]

    return X, y
```

```
def train_val_test_split(X, y, train_ratio=0.8, val_ratio=0.1, test_ratio=0.1):  
    shape=X.shape #Defines the shape of the X matrix  
  
    samples=shape[0]  
  
    X_train=[]  
  
    y_train=[]  
  
    X_val=[] #Setting up empty arrays to later iterate upon  
  
    y_val=[]  
  
    X_test=[]  
  
    y_test=[]  
  
    trainrange=round(samples*train_ratio) #Calculating the range for the training set using the ratio. Using round cause it didn't match with int  
  
    valrange=trainrange+round((val_ratio*samples)) #Same as above  
  
    X_train=X[0:trainrange]  
  
    y_train=y[0:trainrange]  
  
    X_val=X[trainrange:valrange] #Slicing the matrices using the ranges set above  
  
    y_val=y[trainrange:valrange]  
  
    X_test=X[valrange:samples]  
  
    y_test=y[valrange:samples]
```

```
return X_train, y_train, X_val, y_val, X_test, y_test

def scale_features(X_train,X_val,X_test):
    mean=np.mean(X_train, axis=0) #Finding mean + stddev
    std=np.std(X_train, axis=0)
    X_train_scaled=(X_train-mean)/std #Scaling with the same values for each
    X_val_scaled=(X_val-mean)/std
    X_test_scaled=(X_test-mean)/std
    return X_train_scaled,X_val_scaled,X_test_scaled

def calculate_metrics (predicted_labels,true_labels):
    acc=np.mean(predicted_labels==true_labels) #Calculating accuracy as the mean amount of
    labels that equal the true labels
    err=1-acc #Self explanatory
    return acc, err

def knn_single_prediction(new_example,X_train,y_train,k):
    difference=X_train-new_example #Creating a difference matrix to then use linalg.norm on
    dist=[]
```

```
for i in range(len(X_train)):

    dist.append((np.linalg.norm(difference[i]), y_train[i])) #Adding each linalg.norm value for
each element in the array. Had to look up what linalg.norm did

    dist.sort() #Sorting the distances

closest=dist[0:k] #Saying the closest ones the ones with distances from 0 to k

labels=[label for (_, label) in closest] #Had help from groupmembers for this one

values,count=np.unique(labels, return_counts=True) #Finding unique values and counts

max_count=max(count)

for i in range(len(count)):

    if count[i]==max_count: #Saying that if the count is the max count, then the predicted label
is the value at i. Couldn't figure out another way to do this even though its strange

        predicted_label=values[i]

        break

return predicted_label

def predict_labels_knn(X_new,X_train,y_train,k):

predicted_labels=[]

for i in range(len(X_new)):

    new_example=X_new[i]
```

```
label=knn_single_prediction(new_example,X_train,y_train,k) #Iterating the label values
across a new array

predicted_labels.append(label)

return predicted_labels

def tune_k_values(k_values,X_train,y_train,X_val,y_val,is_shuffle):

    metrics = {

        "acc": {"train_acc": [], "val_acc": []},

        "error": {"train_error": [], "val_error": []} #Ripped this straight from the instructions

    }

    best_val=0

    for k in k_values:

        train_label=predict_labels_knn(X_train,X_train,y_train,k) #Getting the labels for each set

        val_label=predict_labels_knn(X_val,X_train,y_train,k)

        train_acc,train_err=calculate_metrics(train_label,y_train) #Calculating error for each set

        val_acc,val_err=calculate_metrics(val_label,y_val)

        metrics["acc"]["train_acc"].append(train_acc)

        metrics["acc"]["val_acc"].append(val_acc) #Appending the metrics

        metrics["error"]["train_error"].append(train_err)
```

```
metrics["error"]["val_error"].append(val_err)

if val_acc>best_val:

    k_good=k #Saying that for each iteration, if the current accuracy is better than the last, then
    the new k value is better

    best_val=val_acc #Same as before but with accuracies

plot_knn_performance(metrics,k_values) #Plotting here cause I already defined the metrics in
this function

return k_good

def plot_knn_performance(metrics,k_values):

    fig,axes = plt.subplots(1, 2, figsize=(15, 6))

    axes[0].plot(k_values, metrics["acc"]["train_acc"], label="Training Accuracy", marker='o')

    axes[0].plot(k_values, metrics["acc"]["val_acc"], label="Validation Accuracy", marker='o')
    #Whole bunch of plot stuff

    axes[0].set_title("Accuracy vs. Number of Neighbors (k)")

    axes[0].set_xlabel("k (Number of Neighbors)")

    axes[0].set_ylabel("Accuracy")

    axes[0].legend()

    axes[0].grid(True)

    axes[1].plot(k_values, metrics["error"]["train_error"], label="Training Error Rate", marker='o')
    #All self explanatory
```

```
axes[1].plot(k_values, metrics["error"]["val_error"], label="Validation Error Rate", marker='o')

axes[1].set_title("Error Rate vs. Number of Neighbors (k)")

axes[1].set_xlabel("k (Number of Neighbors)")

axes[1].set_ylabel("Error Rate")

axes[1].legend()

axes[1].grid(True)

axes[0].axis([0,1000,0,1.0])

axes[1].axis([0,1000,0,1.0])

plt.tight_layout() #To make it cleaner when it pops up

plt.show()

def main():

    path=input("Enter the path to the feature dataset: ") #Importing

    shuffle=input("Shuffle the dataset? (yes/no): ")

    if shuffle == "yes": #Shuffle by using BOO! leans (Cause its halloween)

        shuff=True

    else:

        shuff=False
```

```
seeds=int(input("Enter a seed for loading the dataset: ")) #Defining the seed

features =
["hue_mean","hue_std","saturation_mean","saturation_std","value_mean","value_std","num_lines","has_circle"] #Defining the feature array

label="ClassId" #Setting the label as ClassId

X,y=load_dataset(path,features,label,shuff,seed=seeds) #Loading dataset

X_train,y_train,X_val,y_val,X_test,y_test=train_val_test_split(X,y)

X_train_scaled, X_val_scaled, X_test_scaled = scale_features(X_train, X_val, X_test) #Scaling and getting matrices

print("\nData loaded and split into")

print("\n")

print("Training set\n size:", X_train_scaled.shape[0])

print("Validation set\n size:", X_val_scaled.shape[0]) #Whole lot of printing

print("Test set\n size:", X_test_scaled.shape[0])

k_values=[1, 9, 20, 40, 80, 130, 200, 300, 500, 750, 1000] #Initializing k values

best_k=tune_k_values(k_values, X_train_scaled, y_train, X_val_scaled, y_val, shuff) #Tuning

print(f"\nBased on the plots, the best k appears to be: {best_k:.0f}")

print(f"\nEvaluating final model on test set with k = {best_k:.0f}...") #Same as before

test_label=predict_labels_knn(X_test_scaled, X_train_scaled, y_train, best_k) #Labels

test_acc, test_err = calculate_metrics(test_label, y_test) #Testing metrics
```

```
print("\n--- Final Model Performance ---")

print(f"Test Set Accuracy: {test_acc:.4f}")

print(f"Test Set Error Rate: {test_err:.4f}") #Displaying results.

if __name__ == "__main__":
    main()
```

tp3_team_2_15.py

```
import numpy as np

import math as m

import pandas as pd

import matplotlib.pyplot as plt

from tp3_team_1_15 import load_dataset

from tp3_team_1_15 import train_val_test_split

def calculate_entropy(labels):

    counts = {} #dictionary to count occurrences of labels

    for i in labels:

        counts[i] = counts.get(i, 0) + 1 #adds 1 for each occurrence of i in count
```

```
entropy = 0

for count in counts.values(): #for every element in counts

    prob = count / len(labels) #probability of the label

    entropy += -prob*m.log2(prob) #calculates entropy of that label

return entropy

def information_gain(parent_labels, left_labels, right_labels):

    n_s = len(parent_labels) #total samples

    n_L = len(left_labels) #total left samples

    n_R = len(right_labels) #total right samples

    #calculate entropies

    parent_entropy = calculate_entropy(parent_labels)

    left_entropy = calculate_entropy(left_labels)

    right_entropy = calculate_entropy(right_labels)
```

```
    weighted_entropy = (n_L/n_s)*left_entropy + (n_R/n_s)*right_entropy #calculated weighted entropies
```

```
info_gain = parent_entropy - weighted_entropy #info gain from reducing entropy
```

```
return info_gain
```

```
#function given by task
```

```
def split_dataset(X, y, feature_index, threshold):
```

```
    """
```

Splits a dataset based on a feature and threshold.

```
    """
```

```
X_left, y_left, X_right, y_right = [], [], [], []
```

```
for i, example in enumerate(X):
```

```
    if example[feature_index] < threshold:
```

```
        X_left.append(example)
```

```
        y_left.append(y[i])
```

```
else:  
    X_right.append(example)  
  
    y_right.append(y[i])  
  
return np.array(X_left), np.array(y_left), np.array(X_right), np.array(y_right)  
  
def find_best_split(X, y):  
  
    best_feature = 0 #split parameters  
  
    best_threshold = 0.0  
  
    best_gain = -1  
  
    num_features = X.shape[1] #number of features  
  
    for feat_index in range(num_features): #for every feature  
        features = X[:,feat_index]  
  
        unique_features = np.unique(features) #finds all unique features
```

```
for i in range(len(unique_features) - 1): #for every unique feature

    threshold = (unique_features[i] + unique_features[i+1])/2 #finds threshold via midpoint

    X_left, y_left, X_right, y_right = split_dataset(X, y, feat_index, threshold) #split dataset

    info_gain = information_gain(y, y_left, y_right) #calculate info gained from Y split

    if info_gain > best_gain: #finds split that gives greatest info gain

        best_gain = info_gain

        best_feature = feat_index

        best_threshold = threshold

    return best_feature, best_threshold, best_gain

def build_decision_tree(X, y, max_depth, min_samples_split, current_depth):

    n_s = len(y) #number of samples

    if len(np.unique(y)) == 1: #if theres only 1 unique label aka all same
```

```
predicted_label = y[0]

tree = {"label": predicted_label} #becomes leaf node

return tree

if current_depth >= max_depth or n_s < min_samples_split: #if max depth reached or not enough samples

counts = {} #dictionary to count occurrences of labels

for i in y:

    counts[i] = counts.get(i, 0) + 1 #adds 1 for each occurrence of i in count

common_label = 0

count = 0

for label, count2 in counts.items(): #finds label with max count aka most common

    if count2 > count:

        common_label = label

    count = count2
```

```
tree = {"label": common_label} #becomes leaf node

return tree

best_feature_index, best_threshold_value, best_gain = find_best_split(X, y) #finds conditions
to split

if best_gain <= 0: #if no positive gain

    counts = {} #dictionary to count occurrences of labels

    for i in y:

        counts[i] = counts.get(i, 0) + 1 #adds 1 for each occurrence of i in count

common_label = 0

count = 0

for label, count2 in counts.items(): #iterates through labels and their counts

    if count2 > count: #if current count is greater than previous count, assign new most
common label
```

```
common_label = label
```

```
count = count2
```

```
tree = {"label": common_label} #becomes leaf node
```

```
return tree
```

```
X_left, y_left, X_right, y_right = split_dataset(X, y, best_feature_index, best_threshold_value)  
#splits the dataset
```

```
left_subtree = build_decision_tree(X_left, y_left, max_depth, min_samples_split,  
current_depth + 1) #builds tree for left split
```

```
right_subtree = build_decision_tree(X_right, y_right, max_depth, min_samples_split,  
current_depth + 1) #builds tree for right split
```

```
tree = {
```

```
    "feature": best_feature_index,
```

```
    "threshold": best_threshold_value,
```

```
    "left": left_subtree,
```

```
    "right": right_subtree
```

```
} #saves parameters for each split
```

return tree

#helper

def predict_single(tree, x):

node = tree #top node of the tree

while True: #continues until a leaf

if "label" in node: #if leaf is reached

return node["label"]

feature_index = node["feature"] #get feature index and threshold

threshold = node["threshold"]

if x[feature_index] < threshold: #predict split

node = node["left"]

else:

```
node = node["right"]

def predict(tree, X):
    predicted_labels = [] #array of prediction label
    for feature in X: #loops through rows
        prediction = predict_single(tree, feature) #finds feature label in row
        predicted_labels.append(prediction) #adds to array
    predicted_labels = np.array(predicted_labels) #converts to numpy array
    return predicted_labels

def calculate_accuracy(y_true, y_pred):
    accuracy = sum(y_true == y_pred) / len(y_true) #proportion of correct predictions to total predictions
    return accuracy

#function for plotting data
def plot_dt_performance(train_accs, val_accs, param_values, param_name):
    fig, axes = plt.subplots(1, 1, figsize=(8, 6))
```

```
axes.plot(param_values, train_accs, marker='o', label='Training Accuracy')

axes.plot(param_values, val_accs, marker='o', label='Validation Accuracy')

axes.set_xlabel(param_name)

axes.set_ylabel('Accuracy')

plt.title(f"Decision Tree Accuracy vs {param_name}")

axes.legend()

axes.grid(True)

plt.tight_layout()

plt.show()

def tune_max_depth(X_train, y_train, X_val, y_val):

    depth_values = [1, 2, 3, 4, 5, 7, 10, 15] #range of depth values

    train_accuracies = []

    val_accuracies = []

    for depth in depth_values: #for each depth value

        tree = build_decision_tree(X_train, y_train, depth, min_samples_split=2) #builds tree at depth

        train_accuracies.append(tree.score(X_train, y_train))

        val_accuracies.append(tree.score(X_val, y_val))
```

```
train_predictions = predict(tree, X_train) #makes predictions on training and validation sets

val_predictions = predict(tree, X_val)

train_acc = calculate_accuracy(y_train, train_predictions) #calculate prediction accuracies

val_acc = calculate_accuracy(y_val, val_predictions)

train_accuracies.append(train_acc) #adds accuracies to arrays

val_accuracies.append(val_acc)

plot_dt_performance(train_accuracies, val_accuracies, depth_values, "Max Depth") #plots accuracy

def tune_min_samples_split(X_train, y_train, X_val, y_val, max_depth):

    sample_split_vals = [1, 2, 3, 4, 5, 7, 10, 15] #range of split values

    train_accuracies = []

    val_accuracies = []

    for sample in sample_split_vals:
```

```
tree = build_decision_tree(X_train, y_train, max_depth, sample) #builds tree with split value

train_predictions = predict(tree, X_train) #makes predictions

val_predictions = predict(tree, X_val)

train_acc = calculate_accuracy(y_train, train_predictions) #accuracies of prediction

val_acc = calculate_accuracy(y_val, val_predictions)

train_accuracies.append(train_acc) #stores accuracies

val_accuracies.append(val_acc)

plot_dt_performance(train_accuracies, val_accuracies, sample_split_vals, "Min Samples Split")
#plots accuracies

#function to print tree out

def print_tree(tree, feature_names, indent=""):

    """
    Recursively prints the decision tree in a readable format.
    """

    """
```

```
# Base case: we have reached a leaf node

if "label" in tree:

    # Map the numeric label to a meaningful name

    label_name = "Stop Sign" if tree["label"] == 0 else "Right Turn Sign"

    print(f"{indent}Predict: {label_name}")

return


# Extract information for the current decision node

feature_index = tree["feature"]

feature_name = feature_names[feature_index]

threshold = tree["threshold"]

rule = f"{feature_name} < {threshold:.2f}"


# Recursively print the 'true' branch (left)

print(f"{indent}If {rule}:")

print_tree(tree["left"], feature_names, indent + " ")
```

```
# Recursively print the 'false' branch (right)

print(f"{indent}Else ({rule.replace('<', '>')})")

print_tree(tree["right"], feature_names, indent + " ")

def main():

    filepath = input("Enter the path to the feature dataset: ")

    feature_columns = ['hue_mean', 'hue_std', 'saturation_mean', 'saturation_std', 'value_mean',
    'value_std', 'has_circle', 'num_lines'] #names of all feature columns

    label_column = "ClassId" #name of label col

    X, y = load_dataset(filepath, feature_columns, label_column, shuffle = True, seed=70) #loads
dataset

    X_train, y_train, X_val, y_val, X_test, y_test = train_val_test_split(X, y) #splits datasets

    # prints sizes
```

```
print("\nData loaded and split:")  
  
print(f"Training set: size: {len(X_train)}")  
  
print(f"Validation set: size: {len(X_val)}")  
  
print(f"Test set: size: {len(X_test)}")
```

```
print("\nFinding optimal tree max depth and min samples split...")
```

```
tune_max_depth(X_train, y_train, X_val, y_val) #tests max depths
```

```
optimal_max_depth = int(input("\nEnter the optimal max depth according to your interpretation of the plot:\n"))
```

```
tune_min_samples_split(X_train, y_train, X_val, y_val, optimal_max_depth) #tests sample splits with optimal max_depth
```

```
optimal_min_samples_split = int(input("\nEnter the optimal min samples split according to your interpretation of the plot:\n"))
```

```
print("Evaluating final model on test set with the chosen hyperparameters")
```

```
tree = build_decision_tree(X_train, y_train, optimal_max_depth, optimal_min_samples_split)
#build final tree

#make predictions

y_train_pred = predict(tree, X_train)

y_val_pred = predict(tree, X_val)

y_test_pred = predict(tree, X_test)

#calculate accuracies

train_acc = calculate_accuracy(y_train, y_train_pred)

val_acc = calculate_accuracy(y_val, y_val_pred)

test_acc = calculate_accuracy(y_test, y_test_pred)

#print results

print("\n--- Final Model Performance ---")

print(f"Test Set Accuracy: {test_acc:.4f}")

print(f"Validation Set Accuracy: {val_acc:.4f}")

print(f"Training Set Accuracy: {train_acc:.4f}")
```

```
#print tree results

print("\n--- Final Decision Tree Rules ---")

print_tree(tree, feature_columns)

if __name__ == "__main__":
    main()
```

tp3_team_3_15.py

```
from pathlib import Path

from PIL import Image # Importing libraries

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from tp3_team_1_15 import train_val_test_split, load_dataset #Importing load dataset and
training stuff from task 1.

def calculate_sigmoid(z):

    si=np.clip(z,-400,400) #Clipping the values of Z to prevent overly large values that will mess up
data; range is clipped to (-400,400)
```

```
sigma= 1/(1+np.exp(-si)+(1*10**-15)) #Calculating sigmoid value; epsilon vlaue is added in  
order to prevent division by 0
```

```
return sigma
```

```
def predict_proba(X,w,b):
```

```
z=(np.dot(X,w))+b #Calculating z as the dot product between x and the weight, with the bias  
added
```

```
prob_array = calculate_sigmoid(z)
```

```
return prob_array
```

```
def predict_labels(X,w,b,threshold=0.5):
```

```
prob_array2=predict_proba(X,w,b)
```

```
i=0
```

```
labels=np.zeros(len(prob_array2)) #Forming a zero array to iterate over
```

```
for i in range(len(prob_array2)):
```

```
    if prob_array2[i] >= threshold: #Saying that if the probability value is above the threshold we  
    write it as one, otherwise we write it as 0
```

```
    labels[i]=1
```

```
else:
```

```
    labels[i]=0
```

```
return labels.astype(int) #Returning this array of 1s and 0s

def compute_loss_and_grads(X,y,w,b):

    m=len(X) #Defining the number of values

    prob_array=predict_proba(X,w,b)

    epsilon_prob=prob_array+(1*10**-15) #Adding an epsilon to prevent trying to find the log of 0

    loss=(-1/m)*np.sum(y*np.log(epsilon_prob)+(1-y)*np.log(1-epsilon_prob))

    dw=(1/m)*np.dot(X.T,(prob_array-y)) #Formula given in the instructions; using X.T to find the transpose of X

    db=(1/m)*np.sum((prob_array-y)) #loss,dw,db are calculated as given in the instructions using hte np.sum function

    return loss,dw,db

def train_logistic_regression(X_train, y_train, learn, num):

    shape=X_train.shape #Defining the height of the data

    h=shape[1]

    lost=np.zeros(num) #Making a new empty array

    weight=np.random.randn(h)*0.01 #Defining the initial weight as a random array with the same dimensions as X_train

    b=0 #Setting bias as 0
```



```
return accuracy,error

def main():

    learn=0.4 #Defining the learning rate and number of iterations

    num=2000

    features =
    ["hue_mean","hue_std","saturation_mean","saturation_std","value_mean","value_std","num_li-
nes","has_circle"] #Defining the feature array

    label="ClassId" #Setting the label as ClassId

    X,y=load_dataset("img_features.csv",features,label,shuffle=True,seed=70) #Loading dataset

    X_train,y_train,X_val,y_val,X_test,y_test=train_val_test_split(X,y)

    featmean=np.mean(X_train) #Using the feature mean and standard deviation to standardize
    the data

    featstd=np.std(X_train)

    X_train_std=(X_train-featmean)/featstd

    X_val_std=(X_val-featmean)/featstd #Standardizing as given in the description

    X_test_std=(X_test-featmean)/featstd

    weight,b,lost=train_logistic_regression(X_train_std,y_train,learn,num) #Training

    accval,errval=evaluate_logistic_regression(X_val_std,y_val,weight,b)

    acctest,errtest=evaluate_logistic_regression(X_test_std,y_test,weight,b) #Evaluating the
    validation and test sets
```

```
print(f"Validation accuracy: {accval:.4f}")

print(f"Test accuracy {acctest:.4f}")

print(f"Validation error {errval:.4f}") #Printing the results

print(f"Test error {errest:.4f}")

plt.plot(lost)

plt.plot() #Plotting everything

plt.title("Loss vs Iteration")

plt.xlabel("Iteration")

plt.ylabel("Loss")

plt.show()

if __name__ == "__main__":
    main()
```